

NeuroVisor Architecture Document


Version: Current State **Date:** February 2026 **Author:** Generated from codebase analysis

1. Executive Summary

NeuroVisor is a Firecracker-based microVM orchestrator with integrated LLM agent capabilities. It provides sandboxed code execution environments for AI-driven tasks, combining:

- **Firecracker microVMs** for lightweight, secure isolation
 - **VM Pool** for instant VM availability with pre-warming
 - **gRPC over vsock** for efficient host-guest communication
 - **Ollama integration** for LLM inference with native tool calling
 - **Agent loop** for autonomous code execution tasks
-

2. High-Level Architecture



- node (if installed)
- Timeout enforcement
- stdout/stderr capture

3. Module Breakdown

3.1 Core Modules

Module	Path	Description
vm	src/vm/	Firecracker VM lifecycle management
agent	src/agent/	LLM-driven code execution loop
grpc	src/grpc/	gRPC server/client for host-guest communication
ollama	src/ollama/	Ollama LLM client with tool calling
security	src/security/	Seccomp filters, capabilities, rate limiting
cgroups	src/cgroups/	Resource isolation (CPU/memory limits)
metrics	src/metrics/	Prometheus metrics for observability

3.2 Guest Components

Component	Path	Description
guest_agent	guest/agent/main.rs	In-VM execution server

4. VM Management Subsystem

4.1 VMManger (`src/vm/manager.rs`)

Handles VM creation and destruction:

- Creates Firecracker processes with API sockets


- Configures boot source (kernel, rootfs)
- Sets up vsock devices for host-guest communication
- Assigns unique CIDs (Context IDs) per VM
- Supports snapshot-based boot for faster startup

Configuration:

```
VMManagerConfig {
    kernel_path: "./vmlinuz",
    rootfs_path: "./rootfs.ext4",
    snapshot_path: Option<"./snapshot_file">,
    mem_path: Option<"./mem_file">,
    resource_limits: ResourceLimits,
    vsock_port: 6000,
}
```

4.2 VMPool (`src/vm/pool.rs`)

Thread-safe pool of pre-warmed VMs:



Key Features: - **Pre-warming:** Creates VMs at startup for instant availability -

Acquire/Release: Thread-safe VM checkout/return - **Background Replenisher:**

Maintains target pool size automatically - **Destroy-on-release:** VMs are destroyed after each use for isolation - **Configurable limits:** `warm_size` (default 3), `max_size` (default 10)

4.3 VMHandle (`src/vm/handle.rs`)


Represents a running VM with:

- Unique `vm_id` (UUID v7)
- Firecracker child process
- API socket path
- Vsock socket path
- Context ID (CID)
- Status tracking (Warm/Active)

5. Agent Loop Subsystem

5.1 AgentController (`src/agent/controller.rs`)

Orchestrates LLM-driven code execution:



Configuration:

```
AgentConfig {  
    model: "qwen3",           // LLM model  
    max_iterations: 10,       // Max LLM calls  
    execution_timeout_secs: 30,  
    vsock_port: 6000,  
    connection_retries: 10,  
    connection_retry_delay_ms: 500,  
}
```

5.2 Tool Definition


The agent provides one tool to the LLM:

```
{  
    "type": "function",  
    "function": {  
        "name": "execute_code",  
        "description": "Execute code in a sandboxed environment",  
        "parameters": {  
            "type": "object",  
            "properties": {  
                "language": {  
                    "type": "string",  
                    "enum": ["python", "bash", "javascript"]  
                },  
                "code": {  
                    "type": "string",  
                    "description": "The code to execute"  
                }  
            },  
            "required": ["language", "code"]  
        }  
    }  
}
```

6. Communication Layer

6.1 Vsock Protocol

Firecracker uses vsock for efficient VM communication:



6.2 ExecutionClient ([src/grpc/execution.rs](#))

Host-side gRPC client:

- Performs vsock handshake protocol
- Creates Tonic gRPC channel over vsock
- Provides `execute()` and `execute_with_env()` methods
- Retry logic for guest readiness

6.3 ExecutionServer (Guest)

In-VM gRPC server (`guest/agent/main.rs`):

- Raw libc vsock listener (musl compatible)
- Supports Python, Bash/Shell, JavaScript
- Timeout enforcement
- Streaming output support

Protobuf Definition:

```

service ExecutionService {
    rpc Execute(ExecuteRequest) returns (ExecuteResponse);
    rpc ExecuteStream(ExecuteRequest) returns (stream ExecuteChunk);
}

message ExecuteRequest {
    string language = 1;
    string code = 2;
    uint32 timeout_secs = 3;
    map<string, string> env = 4;
}

message ExecuteResponse {
    string stdout = 1;
    string stderr = 2;
    int32 exit_code = 3;
    double duration_ms = 4;
    bool timed_out = 5;
}

```

7. Ollama Integration

7.1 ChatClient (`src/ollama/tool_use.rs`)

Communicates with Ollama's `/api/chat` endpoint:

- Native tool calling support (model-dependent)
- Fallback text-based tool call parsing
- Temperature=0 for deterministic behavior
- Conversation history management

System Prompt:

You are an AI assistant with the ability to execute code in a sandboxed environment. When you need to run code to accomplish a task, use the `execute_code` tool. Always prefer using bash/shell for simple file operations and system commands.


7.2 Tool Call Parsing

Two parsing modes:

1. **Native:** Model returns `tool_calls` array (qwen3, llama3.2)
2. **Text-based:** Extract JSON tool calls from response text

8. Security Architecture

8.1 Security Layers



8.2 Rate Limiting ([src/security/rate_limit.rs](#))

Token bucket rate limiter for API requests:
- Configurable capacity and refill rate
- Per-request token consumption

9. Observability

9.1 Prometheus Metrics (`src/metrics/mod.rs`)

Inference Metrics: - `neurovisor_requests_total{model}` - Total requests -
`neurovisor_tokens_generated_total{model}` - Tokens generated -
`neurovisor_inference_duration_seconds` - Inference latency -
`neurovisor_errors_total{error_type}` - Error counts

VM Pool Metrics: - `neurovisor_pool_warm_vms` - Pre-warmed VMs -
`neurovisor_pool_active_vms` - Active VMs - `neurovisor_vm_acquire_seconds` - VM acquisition time - `neurovisor_vm_boot_seconds` - VM boot time

Agent Metrics: - `neurovisor_agent_tasks_total{status}` - Agent tasks -
`neurovisor_agent_iterations` - Iterations per task -
`neurovisor_code_execution_seconds` - Code execution time -
`neurovisor_code_executions_total{language,status}` - Executions

Endpoint: `http://0.0.0.0:9090/metrics`

10. File Structure

```
neurovisor/
└── src/
    ├── main.rs          # Daemon entry point
    ├── lib.rs           # Library crate root
    └── agent/
        ├── mod.rs        # Agent module
        └── controller.rs # Agent loop implementation
    └── vm/
        ├── mod.rs        # VM module
        ├── manager.rs    # VMManger
        ├── pool.rs       # VMPool
        ├── handle.rs     # VMHandle
        ├── firecracker.rs # Firecracker API client
        ├── lifecycle.rs   # VM spawn/wait helpers
        └── config.rs      # VM configuration types
    └── grpc/
        ├── mod.rs         # gRPC module
        ├── execution.rs  # ExecutionClient (host-side)
        ├── gateway.rs    # GatewayServer (multi-VM)
        └── server.rs      # InferenceServer
    └── ollama/
        ├── mod.rs         # Ollama module
        ├── client.rs     # OllamaClient (generate)
        └── tool_use.rs    # ChatClient (chat + tools)
    └── security/
        ├── mod.rs         # Security module
        ├── seccomp.rs     # Seccomp BPF filters
        ├── capabilities.rs # Linux capabilities
        └── rate_limit.rs  # Token bucket rate limiter
    └── cgroups/
        ├── mod.rs         # Cgroups module
        └── manager.rs    # CgroupManager
    └── metrics/
        └── mod.rs         # Prometheus metrics
    └── guest/
        └── agent/
            └── main.rs      # Guest execution server
    └── proto/
        ├── execution.proto # Execution service proto
        └── inference.proto # Inference service proto
    └── vmlinuz           # Linux kernel
    └── rootfs.ext4        # Alpine Linux rootfs
    └── firecracker        # Firecracker binary (symlink)
```

11. Configuration

11.1 Command Line Arguments

```
# Default daemon mode  
sudo ./neurovisor  
  
# Custom pool size  
sudo ./neurovisor --warm 5 --max 20  
  
# Use snapshots for faster boot  
sudo ./neurovisor --snapshot  
  
# Agent mode (single task)  
sudo ./neurovisor --agent "Write a script that prints hello world"
```

11.2 Constants

Constant	Value	Description
KERNEL_PATH	./vmlinuz	Linux kernel path
ROOTFS_PATH	./rootfs.ext4	Root filesystem path
METRICS_PORT	9090	Prometheus metrics port
GATEWAY_PORT	50051	gRPC gateway port
VSOCK_PORT	6000	Guest execution service port
DEFAULT_WARM_SIZE	3	Default pre-warmed VMs
DEFAULT_MAX_SIZE	10	Maximum VMs

12. Data Flow Examples

12.1 Agent Task Execution

1. User runs: ./neurovisor --agent "List files in /tmp"
2. Main → AgentController.run("List files in /tmp")
3. AgentController → Ollama /api/chat
 - System prompt + user task + tools
4. Ollama returns tool_call:

```
{  "name": "execute_code",  "arguments": {"language": "bash", "code": "ls -la /tmp"}}
```
5. AgentController.execute_code():
 - a. pool.acquire() → VMHandle
 - b. ExecutionClient.connect(vsock_path)
 - c. Vsock handshake (CONNECT/OK)
 - d. gRPC Execute(language="bash", code="ls -la /tmp")
 - e. Guest runs /bin/sh -c "ls -la /tmp"
 - f. Returns stdout/stderr/exit_code
 - g. pool.release(vm) → VM destroyed
6. AgentController → Ollama /api/chat
 - Add tool result to conversation
7. Ollama returns final response (no tool calls)
8. AgentController returns AgentResult

12.2 VM Lifecycle

1. `VMPool.initialize()`
 - └─ Creates 3 warm VMs
2. `VMPool.acquire()`
 - └─ Returns pre-warmed VM, marks as active
3. [Code execution happens]
4. `VMPool.release(vm)`
 - └─ Destroys VM (kill Firecracker, cleanup)
 - └─ Triggers replenisher
5. Background replenisher
 - └─ Creates new warm VM to maintain pool size

13. Dependencies

13.1 Rust Crates

Crate	Purpose
<code>tokio</code>	Async runtime
<code>tonic</code>	gRPC framework
<code>prost</code>	Protobuf codegen
<code>hyper</code>	HTTP server
<code>reqwest</code>	HTTP client (Ollama)
<code>serde_json</code>	JSON serialization
<code>prometheus</code>	Metrics
<code>uuid</code>	UUID generation
<code>lazy_static</code>	Static initialization

13.2 External Services

Service	Purpose	Default URL
Ollama	LLM inference	<code>http://localhost:11434</code>
Firecracker	MicroVM hypervisor	<code>./firecracker</code> binary

14. Future Considerations

- Networking:** Add VM network support for internet access
- Persistence:** Support persistent VM state across requests
- Multi-tenant:** Add user/namespace isolation
- Streaming:** Full streaming support for long-running tasks
- Languages:** Add more language runtimes to guest rootfs
- GPU:** GPU passthrough for ML workloads

Document generated from NeuroVisor codebase analysis - February 2026